

# **Swin Transformer:** **Hierarchical Vision Transformer** **using Shifted Windows**

---

Microsoft Research Asia  
2021

# 1. Introduction

---

- ✓ Transformer is notable for its use of attention to model long-range dependencies in the data.
- ✓ In this paper, we seek to expand the applicability of Transformer such that it can serve as a general-purpose backbone for computer vision.
- ✓ **There are two challenges:**
  - Differences of scale(variability in scale)
  - Higher resolution of pixels in images compared to words in passages of text
- ✓ To overcome these issues, we propose a general-purpose Transformer backbone **"Swin Transformer"**

## 2. Related Work

---

- ✓ CNN and variants
  - Served as the standard network model throughout computer vision.
  - VGG, GoogleNet, ResNet, DenseNet ...
- ✓ Self-attention based backbone architectures
  - some works employ self-attention layers to replace some or all of the spatial convolution layers in the popular ResNet.
- ✓ Self-attention/Transformers to complement CNNs
- ✓ Transformer based vision backbones
  - **Vision Transformer(ViT)**
  - An image is 11 worth 16x16 words: Transformers for image recognition at scale(2021)

# 3. Method

## ❖ Structure

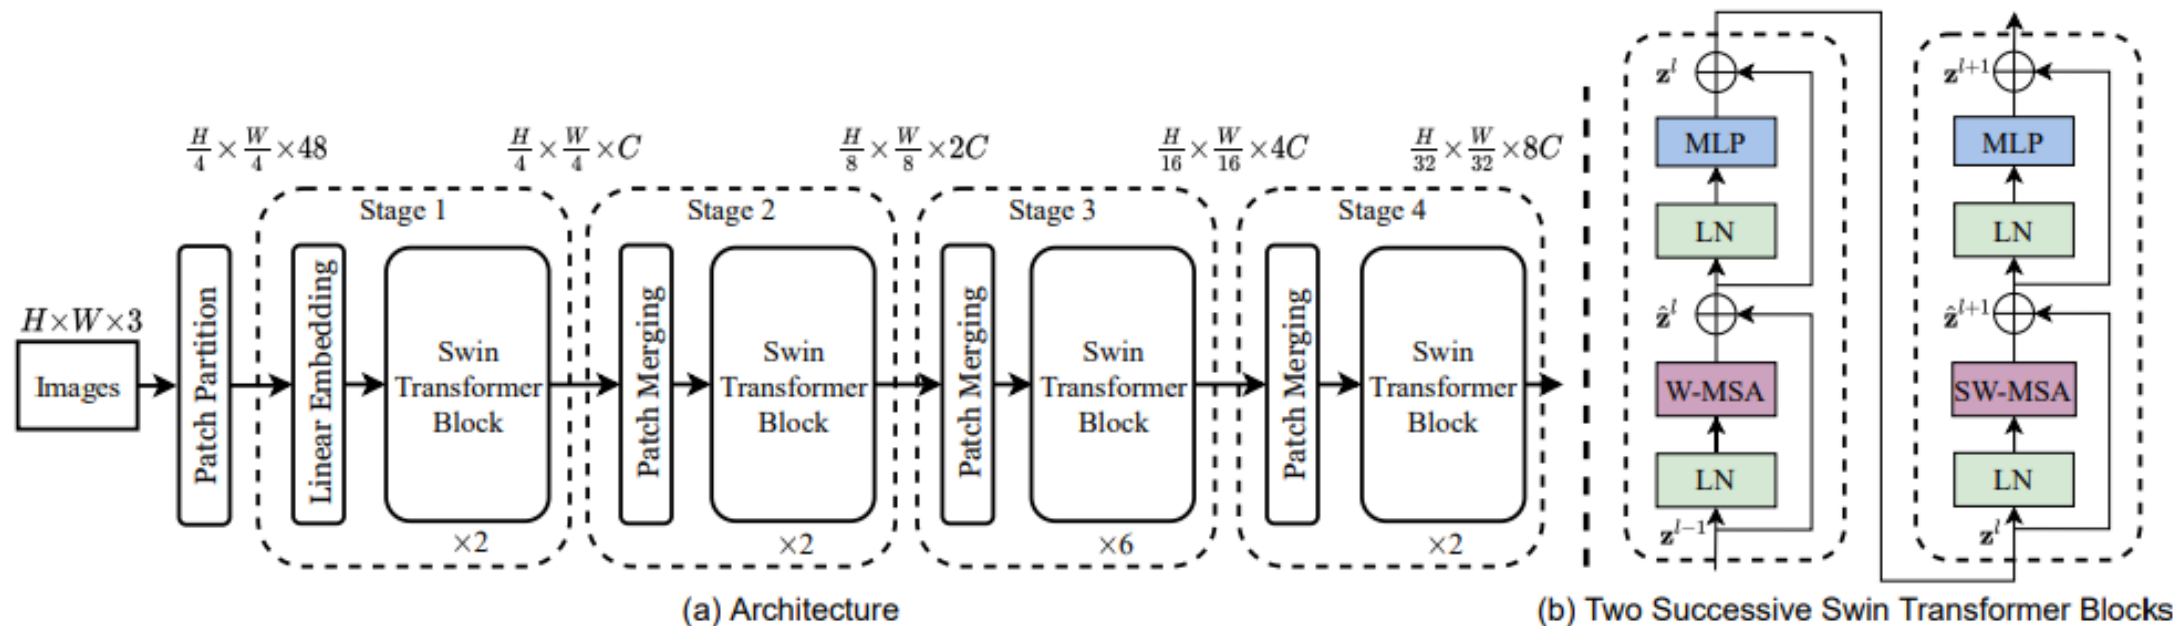



Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

# 3. Method


## ❖ Patch Partition + Linear Embedding



- ✓ It first splits an input RGB image into non-overlapping, like ViT.
- ✓ Used a patch size of  $4 \times 4$ (in Tiny).
  - Thus the feature dimension of each patch is  $4 \times 4 \times 3 = 48$
- ✓ A linear embedding layer is applied on this raw-valued feature to project it to an arbitrary dimension.

# 3. Method

## ❖ Patch Merging




- ✓ The first patch merging layer concatenates the features of each group of  $2 \times 2$  neighboring patches and applies a linear layer on the  $4C$ -dimensional concatenated features.
- ✓ This reduces the number of tokens by a multiple of  $2 \times 2 = 4$  ( $2 \times$  downsampling of resolution), and the output dimension is set to  $2C$ .
- ✓ To produce a hierarchical representation, the number of tokens is reduced by patch merging layers as the network gets deeper.

# 3. Method

## Swin Transformer

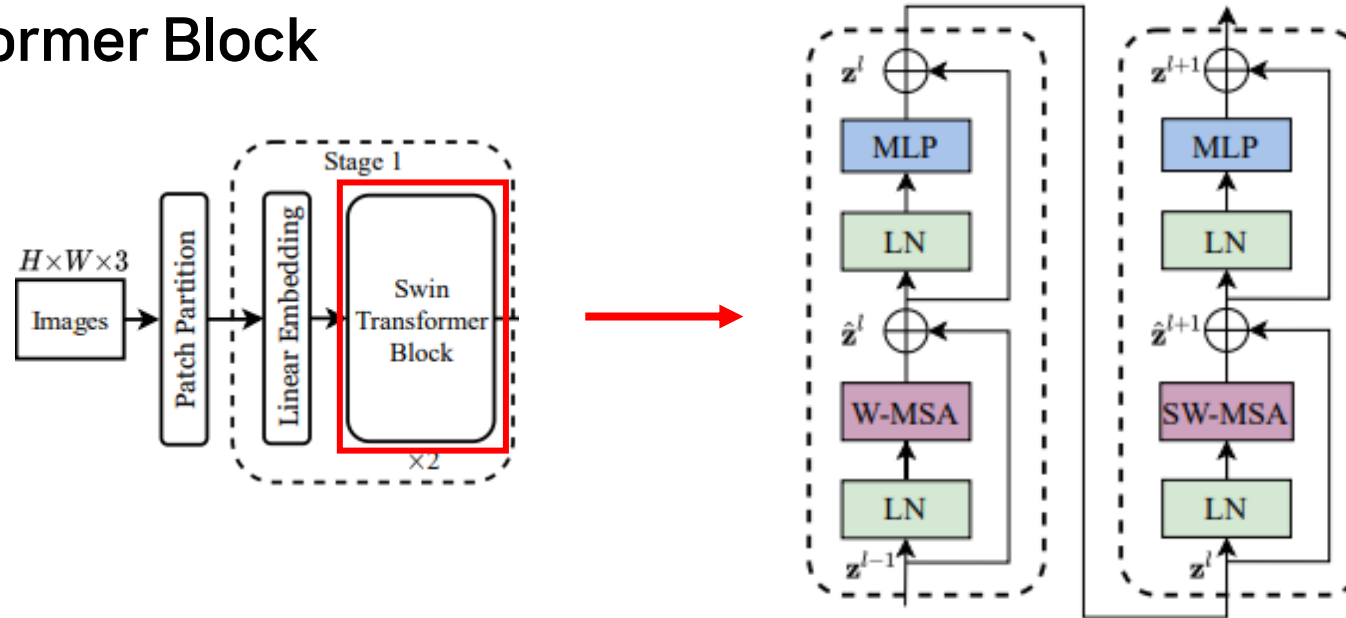
### Patch Merging



CSDN @太阳花的小绿豆

# 3. Method

## ❖ Swin Transformer Block



$$\hat{\mathbf{z}}^l = \text{W-MSA}(\text{LN}(\mathbf{z}^{l-1})) + \mathbf{z}^{l-1},$$

$$\mathbf{z}^l = \text{MLP}(\text{LN}(\hat{\mathbf{z}}^l)) + \hat{\mathbf{z}}^l,$$

$$\hat{\mathbf{z}}^{l+1} = \text{SW-MSA}(\text{LN}(\mathbf{z}^l)) + \mathbf{z}^l,$$



$$\mathbf{z}^{l+1} = \text{MLP}(\text{LN}(\hat{\mathbf{z}}^{l+1})) + \hat{\mathbf{z}}^{l+1},$$



# 3. Method

## ❖ W-MSA & SW-MSA


- ✓ W-MSA: Local Window 안에서 self attention
- ✓ SW-MSA: Local Window 간의 self attention





# 3. Method

## ❖ Efficient batch computation

$$x \in \mathbb{R}^{B \times N_h \times N_w \times C} \quad x \in \mathbb{R}^{nB \times M \times M \times C}, n = \frac{N_h}{M} + \frac{N_w}{M}$$



$$x \in \mathbb{R}^{1 \times 8 \times 8 \times 3}$$




$$x \in \mathbb{R}^{4 \times 4 \times 4 \times 3}$$

$$M = 4, n = 4$$

$$(batch * patch_h * patch_w * channel)$$

$$\left( batch' * \frac{patch_h}{M} * \frac{patch_w}{M} * channel \right)$$

$$batch' = batch * \frac{patch_h}{m} * \frac{patch_w}{M}$$




# 3. Method

## ❖ SW-MSA


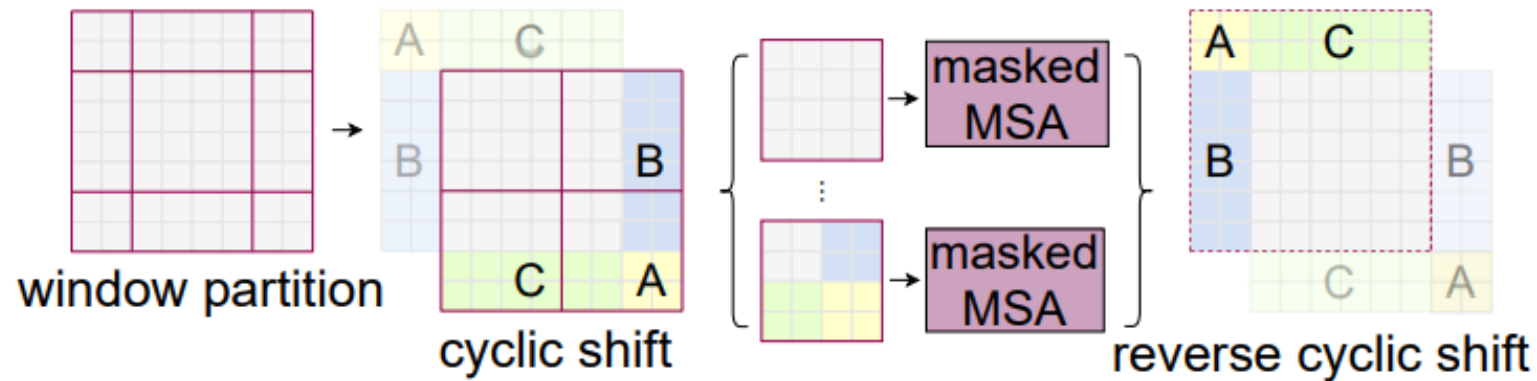
- ✓ Local Window 간의 self attention
- ✓ SW-MSA 수행 시에 window 개수가 H, W 별로 1개 씩 늘어나게 됨
  - 더 많은 window 사용으로 인한 비효율적인 연산 수행

## ❖ Cyclic Shift + Attention Mask를 통한 W-MSA와 동일한 window 개수 사용



# 3. Method



## ❖ Cyclic shift & Attention Mask



- ✓ SW-MSA를 수행하여 나뉘어진 window로 구성되어야 했던 것을 일정 크기 만큼 shift한다.
  - $Shift\ size = Window\ size(M) // 2$
- ✓ 색이 다른 부분 모두 각각 attention이 적용되어야 하기 때문에 mask를 적용한다.
  - ✓ Why? 실제로 색이 다른 부분은 이미지 상에서 인접한 부분이 아니기 때문이다.
- ✓ W-MSA와 동일하게 4개의 window만 사용하여 local window간에 attention을 계산한다.
- ✓ 계산 결과에 다시 shift를 적용하여 결과를 복원한다.

# 3. Method

## ❖ Cyclic shift & Attention Mask



<정리>

- ✓ Shift 후 W-MSA처럼 4개의 Window만 사용하여 Window 간의 Self Attention 수행
- ✓ 이후 Attention Mask를 씌워 Mask에서 따로 Attention을 적용한다.
  - Relative bias 적용 후 Attention Mask가 적용됨

# 3. Method

## ❖ Relative position bias

$$\text{Attention}(Q, K, V) = \text{SoftMax}(QK^T / \sqrt{d} + B)V,$$

- ✓ In computing self-attention, we apply a relative position bias to each head in computing similarity
  - $Q, K, V \in R^{M^2 \times d}$ : *query*, *key* and *value* matrices
  - $d$  = query/key dimension
  - $M^2$ : the number of patches in a window

# 3. Method

## ❖ Relative position bias

$$\text{Attention}(Q, K, V) = \text{SoftMax}(QK^T / \sqrt{d} + B)V,$$

- ✓ 두 축마다 Relative position의 범위:  $[-M + 1, M - 1]$
- ✓ Bias Index Matrix:  $\hat{B} \in R^{(2M-1) \times (2M-1)}$
- ✓  $B$ 는  $\hat{B}$ 의 값을 사용

# 3. Method

## ❖ Relative position bias

<i>x axis</i>			1	2	3	4
1	2	1	0	0	-1	-1
3	4	2	0	0	-1	-1
		3	1	1	0	0
		4	1	1	0	0

<i>y axis</i>			1	2	3	4
1	2	1	0	-1	0	-1
3	4	2	1	0	1	0
		3	0	-1	0	-1
		4	1	0	1	0

- ✓ 축 별로 상대적인 거리를 계산한다.
  - 1과 같은 축에 있는 값들은 0, 1칸 차이냐면 1 or -1, 2칸 차이냐면 - or -2
- ✓ 이렇게 구한 각 Matrix에 (window size - 1)값을 더해 준다.
  - 실제로 index로 나타내기 위해 범위가 0부터 시작되도록 변환하기 위함



# 3. Method

## ❖ Relative position bias

*Relative Position Index*

	1	2	3	4	5	6	7	8	9
1	12	11	10	7	6	5	2	1	0
2	13	12	11	8	7	6	3	2	1
3	14	13	12	9	8	7	4	3	2
4	17	16	15	12	11	10	7	6	5
5	18	17	16	13	12	11	8	7	6
6	19	18	17	14	13	12	9	8	7
7	22	21	20	17	16	15	12	11	10
8	23	22	21	18	17	16	13	12	11
9	24	23	22	19	18	17	14	13	12


- ✓  $x\_axis\_matrix \ast = 2 \ast window\_size - 1$
- ✓  $relative\_position\_matrix = x\_axis\_matrix + y\_axis\_matrix$
- ✓ 이 행렬의 크기는  $(2M - 1) \ast (2M - 1)$

# 3. Method

## ❖ Relative position bias

*Relative Position Index*

	1	2	3	4	5	6	7	8	9
1	12	11	10	7	6	5	2	1	0
2	13	12	11	8	7	6	3	2	1
3	14	13	12	9	8	7	4	3	2
4	17	16	15	12	11	10	7	6	5
5	18	17	16	13	12	11	8	7	6
6	19	18	17	14	13	12	9	8	7
7	22	21	20	17	16	15	12	11	10
8	23	22	21	18	17	16	13	12	11
9	24	23	22	19	18	17	14	13	12



- ✓ 이렇게 만든 Relative position index를  $\hat{B}$ 에서 값을 조회하여  $B$ 라는 행렬을 구성하게 된다.
- ✓  $B$ 를 Attention 수식에 적용한다.

# 4. Experiments

(a) Regular <u>ImageNet-1K</u> trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
RegNetY-4G [48]	224 <sup>2</sup>	21M	4.0G	1156.7	80.0
RegNetY-8G [48]	224 <sup>2</sup>	39M	8.0G	591.6	81.7
RegNetY-16G [48]	224 <sup>2</sup>	84M	16.0G	334.7	82.9
EffNet-B3 [58]	300 <sup>2</sup>	12M	1.8G	732.1	81.6
EffNet-B4 [58]	380 <sup>2</sup>	19M	4.2G	349.4	82.9
EffNet-B5 [58]	456 <sup>2</sup>	30M	9.9G	169.1	83.6
EffNet-B6 [58]	528 <sup>2</sup>	43M	19.0G	96.9	84.0
EffNet-B7 [58]	600 <sup>2</sup>	66M	37.0G	55.1	84.3
ViT-B/16 [20]	384 <sup>2</sup>	86M	55.4G	85.9	77.9
ViT-L/16 [20]	384 <sup>2</sup>	307M	190.7G	27.3	76.5
DeiT-S [63]	224 <sup>2</sup>	22M	4.6G	940.4	79.8
DeiT-B [63]	224 <sup>2</sup>	86M	17.5G	292.3	81.8
DeiT-B [63]	384 <sup>2</sup>	86M	55.4G	85.9	83.1
Swin-T	224 <sup>2</sup>	29M	4.5G	755.2	81.3
Swin-S	224 <sup>2</sup>	50M	8.7G	436.9	83.0
Swin-B	224 <sup>2</sup>	88M	15.4G	278.1	83.5
Swin-B	384 <sup>2</sup>	88M	47.0G	84.7	84.5

(b) <u>ImageNet-22K</u> pre-trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
R-101x3 [38]	384 <sup>2</sup>	388M	204.6G	-	84.4
R-152x4 [38]	480 <sup>2</sup>	937M	840.5G	-	85.4
ViT-B/16 [20]	384 <sup>2</sup>	86M	55.4G	85.9	84.0
ViT-L/16 [20]	384 <sup>2</sup>	307M	190.7G	27.3	85.2
Swin-B	224 <sup>2</sup>	88M	15.4G	278.1	85.2
Swin-B	384 <sup>2</sup>	88M	47.0G	84.7	86.4
Swin-L	384 <sup>2</sup>	197M	103.9G	42.1	87.3

Table 1. Comparison of different backbones on ImageNet-1K classification. Throughput is measured using the GitHub repository of [68] and a V100 GPU, following [63].

# 4. Experiments

(a) Various frameworks							
Method	Backbone	AP <sup>box</sup>	AP <sup>box</sup> <sub>50</sub>	AP <sup>box</sup> <sub>75</sub>	#param.	FLOPs	FPS
Cascade	R-50	46.3	64.3	50.5	82M	739G	18.0
Mask R-CNN	Swin-T	<b>50.5</b>	<b>69.3</b>	<b>54.9</b>	86M	745G	15.3
ATSS	R-50	43.5	61.9	47.0	32M	205G	28.3
	Swin-T	<b>47.2</b>	<b>66.5</b>	<b>51.3</b>	36M	215G	22.3
RepPointsV2	R-50	46.5	64.6	50.3	42M	274G	13.6
	Swin-T	<b>50.0</b>	<b>68.5</b>	<b>54.2</b>	45M	283G	12.0
Sparse R-CNN	R-50	44.5	63.4	48.2	106M	166G	21.0
	Swin-T	<b>47.9</b>	<b>67.3</b>	<b>52.3</b>	110M	172G	18.4

(b) Various backbones w. Cascade Mask R-CNN									
	AP <sup>box</sup>	AP <sup>box</sup> <sub>50</sub>	AP <sup>box</sup> <sub>75</sub>	AP <sup>mask</sup>	AP <sup>mask</sup> <sub>50</sub>	AP <sup>mask</sup> <sub>75</sub>	#param	FLOPs	FPS
DeiT-S <sup>†</sup>	48.0	67.2	51.7	41.4	64.2	44.3	80M	889G	10.4
R50	46.3	64.3	50.5	40.1	61.7	43.4	82M	739G	18.0
Swin-T	<b>50.5</b>	<b>69.3</b>	<b>54.9</b>	<b>43.7</b>	<b>66.6</b>	<b>47.1</b>	86M	745G	15.3
X101-32	48.1	66.5	52.4	41.6	63.9	45.2	101M	819G	12.8
Swin-S	<b>51.8</b>	<b>70.4</b>	<b>56.3</b>	<b>44.7</b>	<b>67.9</b>	<b>48.5</b>	107M	838G	12.0
X101-64	48.3	66.4	52.3	41.7	64.0	45.1	140M	972G	10.4
Swin-B	<b>51.9</b>	<b>70.9</b>	<b>56.5</b>	<b>45.0</b>	<b>68.4</b>	<b>48.7</b>	145M	982G	11.6

(c) System-level Comparison						
Method	mini-val		test-dev		#param. FLOPs	
	AP <sup>box</sup>	AP <sup>mask</sup>	AP <sup>box</sup>	AP <sup>mask</sup>		
RepPointsV2* [12]	-	-	52.1	-	-	-
GCNet* [7]	51.8	44.7	52.3	45.4	-	1041G
RelationNet++* [13]	-	-	52.7	-	-	-
SpineNet-190 [21]	52.6	-	52.8	-	164M	1885G
ResNeSt-200* [78]	52.5	-	53.3	47.1	-	-
EfficientDet-D7 [59]	54.4	-	55.1	-	77M	410G
DetectoRS* [46]	-	-	55.7	<b>48.5</b>	-	-
YOLOv4 P7* [4]	-	-	55.8	-	-	-
Copy-paste [26]	55.9	47.2	<b>56.0</b>	47.4	185M	1440G
X101-64 (HTC++)	52.3	46.0	-	-	155M	1033G
Swin-B (HTC++)	56.4	49.1	-	-	160M	1043G
Swin-L (HTC++)	57.1	49.5	57.7	50.2	284M	1470G
Swin-L (HTC++)*	<b>58.0</b>	<b>50.4</b>	<b>58.7</b>	<b>51.1</b>	284M	-

Table 2. Results on COCO object detection and instance segmentation. <sup>†</sup> denotes that additional decovolution layers are used to produce hierarchical feature maps. \* indicates multi-scale testing.

# 4. Experiments

ADE20K		val	test	#param.	FLOPs	FPS
Method	Backbone	mIoU	score			
DANet [23]	ResNet-101	45.2	-	69M	1119G	15.2
DLab.v3+ [11]	ResNet-101	44.1	-	63M	1021G	16.0
ACNet [24]	ResNet-101	45.9	38.5	-		
DNL [71]	ResNet-101	46.0	56.2	69M	1249G	14.8
OCRNet [73]	ResNet-101	45.3	56.0	56M	923G	19.3
UperNet [69]	ResNet-101	44.9	-	86M	1029G	20.1
OCRNet [73]	HRNet-w48	45.7	-	71M	664G	12.5
DLab.v3+ [11]	ResNeSt-101	46.9	55.1	66M	1051G	11.9
DLab.v3+ [11]	ResNeSt-200	48.4	-	88M	1381G	8.1
SETR [81]	T-Large <sup>‡</sup>	50.3	61.7	308M	-	-
UperNet	DeiT-S <sup>†</sup>	44.0	-	52M	1099G	16.2
UperNet	Swin-T	46.1	-	60M	945G	18.5
UperNet	Swin-S	49.3	-	81M	1038G	15.2
UperNet	Swin-B <sup>‡</sup>	51.6	-	121M	1841G	8.7
UperNet	Swin-L <sup>‡</sup>	53.5	62.8	234M	3230G	6.2

Table 3. Results of semantic segmentation on the ADE20K val and test set. <sup>†</sup> indicates additional deconvolution layers are used to produce hierarchical feature maps. <sup>‡</sup> indicates that the model is pre-trained on ImageNet-22K.

## 5. Ablation Study

	ImageNet		COCO		ADE20k
	top-1	top-5	AP <sup>box</sup>	AP <sup>mask</sup>	mIoU
w/o shifting	80.2	95.1	47.7	41.5	43.3
shifted windows	<b>81.3</b>	<b>95.6</b>	<b>50.5</b>	<b>43.7</b>	<b>46.1</b>
no pos.	80.1	94.9	49.2	42.6	43.8
abs. pos.	80.5	95.2	49.0	42.4	43.2
abs.+rel. pos.	81.3	95.6	50.2	43.4	44.0
rel. pos. w/o app.	79.3	94.7	48.2	41.9	44.1
rel. pos.	<b>81.3</b>	<b>95.6</b>	<b>50.5</b>	<b>43.7</b>	<b>46.1</b>

Table 4. Ablation study on the *shifted windows* approach and different position embedding methods on three benchmarks, using the Swin-T architecture. w/o shifting: all self-attention modules adopt regular window partitioning, without *shifting*; abs. pos.: absolute position embedding term of ViT; rel. pos.: the default settings with an additional relative position bias term (see Eq. (4)); app.: the first scaled dot-product term in Eq. (4).

## 5. Ablation Study

method	MSA in a stage (ms)				Arch. (FPS)		
	S1	S2	S3	S4	T	S	B
sliding window (naive)	122.5	38.3	12.1	7.6	183	109	77
sliding window (kernel)	7.6	4.7	2.7	1.8	488	283	187
Performer [14]	4.8	2.8	1.8	1.5	638	370	241
window (w/o shifting)	2.8	1.7	1.2	0.9	770	444	280
shifted window (padding)	3.3	2.3	1.9	2.2	670	371	236
shifted window (cyclic)	3.0	1.9	1.3	1.0	755	437	278

Table 5. Real speed of different self-attention computation methods and implementations on a V100 GPU.

	Backbone	ImageNet		COCO		ADE20k
		top-1	top-5	AP <sup>box</sup>	AP <sup>mask</sup>	
sliding window	Swin-T	81.4	95.6	50.2	43.5	45.8
Performer [14]	Swin-T	79.0	94.2	-	-	-
shifted window	Swin-T	81.3	95.6	50.5	43.7	46.1

Table 6. Accuracy of Swin Transformer using different methods for self-attention computation on three benchmarks.

## 5. Conclusion

---

- ✓ This paper presents Swin Transformer, a new vision Transformer which produces a hierarchical feature representation.
- ✓ Swin Transformer achieves the state-of-the-art performance on COCO object detection and ADE20K semantic segmentation
- ✓ As a key element of Swin Transformer, the shifted window based self-attention is shown to be effective and efficient on vision problems, and we look forward to investigating its use in natural language processing as well.